

## DS8 - DBSCAN assignment

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import NearestNeighbors

n1 = 100
X1_x = (np.random.rand(n1)-2)*2
X1_y = np.random.rand(n1)+1

n2 = 100
X2_x = (np.random.rand(n2)+1)*2
X2_y = np.random.rand(n2)+1

step = 0.01
X3_x = np.arange(-2, 2+step, step)
n3 = len(X3_x)
X3_y = (np.random.randn())/3-0.001)*X3_x**2 - 2 + np.random.randn(n3)/3

X1 = np.vstack([X1_x, X1_y])
X2 = np.vstack([X2_x, X2_y])
X3 = np.vstack([X3_x, X3_y])
X = np.hstack([X1, X2, X3]).T
X[:5, :]
```

```
array([[ -3.16896671,  1.10362081],
       [ -3.56468624,  1.14737778],
       [ -3.89786797,  1.00955461],
       [ -3.16509829,  1.32067596],
       [ -3.66480546,  1.29687086]])
```

1. Determine good parameters to cluster  $X$  with DBSCAN. Perform DBSCAN (use sklearn implementation) on the data and visualise the plot

```
X_scaled = StandardScaler().fit_transform(X)

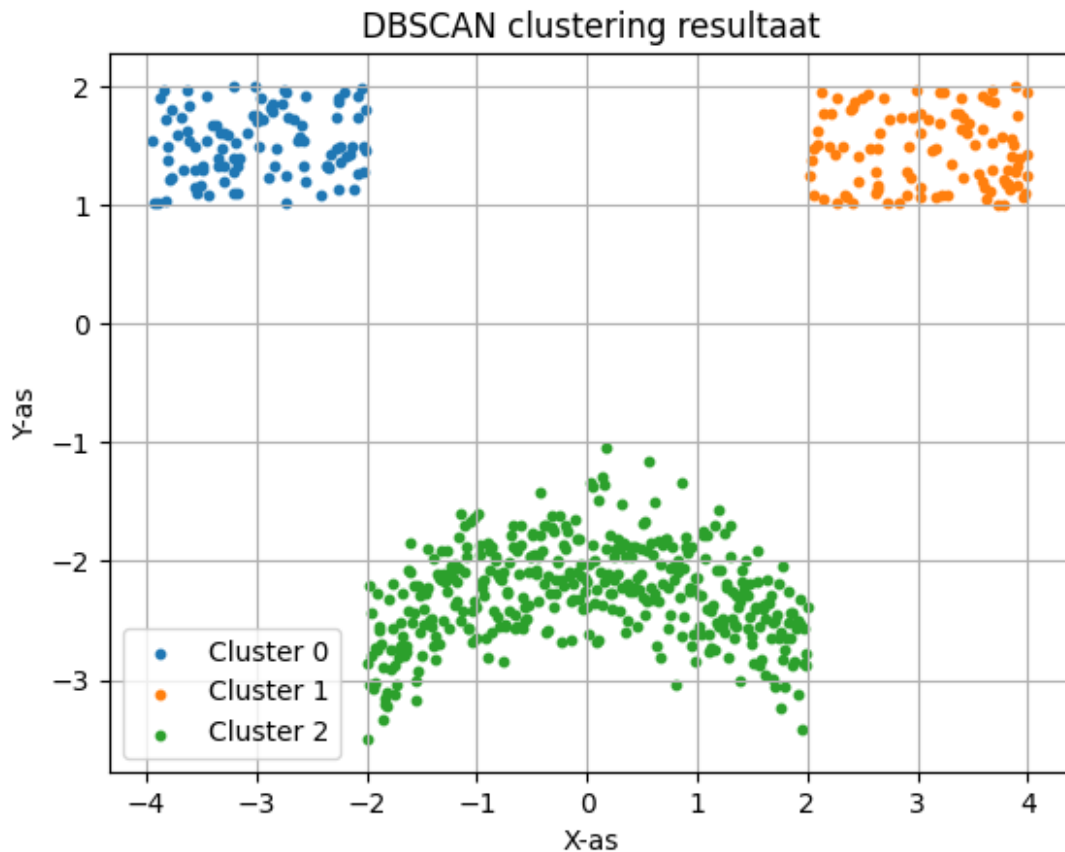
from sklearn.cluster import DBSCAN
dbscan = DBSCAN(eps=0.3, min_samples=5)
labels = dbscan.fit_predict(X_scaled)

unique_labels = set(labels)

for label in unique_labels:
    mask = (labels == label)
```

```
plt.scatter(X[mask, 0], X[mask, 1], s=10, label=f'Cluster {label}')
if label != -1 else 'Noise')

plt.title('DBSCAN clustering resultaat')
plt.xlabel('X-as')
plt.ylabel('Y-as')
plt.legend()
plt.grid(True)
plt.show()
```



## 2. Write a DBSCAN implementation that clusters data from arbitrary dimension

For the following exercises, no for-loops are allowed. Unless stated otherwise.

We define a couple of constants (recognizable by writing them in ALL CAPS). You can use these constants in- and outside of the function.

```
EPSILON = 0.7 # our value for epsilon
INDEX = np.arange(len(X)) # an array of all the indexes of X
NOISE = -1 # which value is used as the noise cluster
MIN_NEIGHBOURS = 4 # our value for n_min
```

We need to compute the distance between all the datapoints and a specific datapoint very often. That's why it is important to write a very fast distance function.

2.1 Make a function `dist(data: np.ndarray, p_index: int) -> np.ndarray` that outputs the Euclidean distance between all data points in  $X$  and point  $p = X_{p\_index}$

```
def dist(data: np.ndarray, p_index: int) -> np.ndarray:
    p = data[p_index]
    diff = data - p
    return np.sqrt(np.sum(diff**2, axis=1))
```

Test the speed of your code with the following code block. Your function should run in less than 50µs.

```
%%timeit
# test your distance function
dist(X, 0)
```

8.32 µs ± 109 ns per loop (mean ± std. dev. of 7 runs, 100,000 loops each)

We also want to extract all the candidate neighbours for a certain point  $p$ . Remember, if a datapoint is already assigned to a cluster, it is not a neighbour anymore. That's why we are looking for candidates.

2.2 Make a function `determine_neighbours(data: np.ndarray, p_index: int) -> candidates: np.ndarray` that outputs a boolean array `candidates` where `candidatesi = True` if point with index  $i$  is a potential neighbour of point  $p$ .

Also make sure that point  $p$  is not assigned as a candidate neighbour of itself.

```
def determine_neighbours(data: np.ndarray, p_index: int) ->
np.ndarray:
    distances = dist(data, p_index)
    candidates = (distances <= EPSILON)
    candidates[p_index] = False
    return candidates
```

```
%%timeit
determine_neighbours(X, 0)
```

9.73 µs ± 93.7 ns per loop (mean ± std. dev. of 7 runs, 100,000 loops each)

Your function should run in less than 100µs.

2.3 Make a function `find_neighbours(clusters: np.ndarray, data: np.ndarray, p_index: int) -> neighbours: np.ndarray` that outputs the neighbours of point  $p$ , where `clusters` is the array with cluster labels for each point.

```
def find_neighbours(clusters: np.ndarray, data: np.ndarray, p_index:
int) -> np.ndarray:
    distances = dist(data, p_index)
    unassigned = (clusters == NOISE)
    close_enough = (distances <= EPSILON)
    not_self = np.arange(len(data)) != p_index
    neighbours = unassigned & close_enough & not_self
    return neighbours
```

```
%%timeit
find_neighbours(np.zeros(len(X)), X, 0)
```

15.7  $\mu$ s  $\pm$  136 ns per loop (mean  $\pm$  std. dev. of 7 runs, 100,000 loops each)

Your function should run in less than 150 $\mu$ s.

Now we are going to write the core part of the DBSCAN algorithm. You can use at most 1 for-loop for this function.

2.4 Make a function `assign_neighbours_to_cluster(clusters: np.ndarray, data: np.ndarray, p_index: int, c: int)` that modifies the array `clusters`.

Make sure the function does the following steps

1. Check if `p_index` is already assigned to a cluster, if so return `None`
2. If not, change the cluster value of point `p_index` to `c`
3. Check if point  $p$  has enough neighbours to assign those neighbours as well to cluster `c`
4. Recursively add the neighbours and their neighbours of point  $p$  to cluster `c`.

```
def assign_neighbours_to_cluster(clusters: np.ndarray, data:
np.ndarray, p_index: int, c: int):
    if clusters[p_index] != NOISE:
        return None

    clusters[p_index] = c

    neighbours = find_neighbours(clusters, data, p_index)
    if np.sum(neighbours) < MIN_NEIGHBOURS:
        return None
```

```

clusters[neighbours] = c
to_check = np.where(neighbours)[0]

i = 0
while i < len(to_check):
    point = to_check[i]
    if clusters[point] == NOISE:
        clusters[point] = c
    new_neigh = find_neighbours(clusters, data, point)
    if np.sum(new_neigh) >= MIN_NEIGHBOURS:
        new_points = np.where(new_neigh & (clusters == NOISE))[0]
        clusters[new_points] = c
        to_check = np.concatenate([to_check, new_points])
    i += 1

%%timeit
assign_neighbours_to_cluster(np.zeros(len(X)), X, 0, 1)

547 ns ± 2.84 ns per loop (mean ± std. dev. of 7 runs, 1,000,000 loops
each)

```

Your function should run in less than 1ms.

## 2.5 Combine all these elements into a DBSCAN algorithm. You are allowed to use 1 for loop.

Compare your results from 1.

```

def DBSCAN(data: np.ndarray) -> np.ndarray:
    clusters = np.full(len(data), NOISE) # start: alles is noise
    current_cluster = 0

    for p_index in range(len(data)):
        if clusters[p_index] != NOISE:
            continue
        # Probeer een nieuw cluster te maken vanuit dit punt
        previous = clusters.copy()
        assign_neighbours_to_cluster(clusters, data, p_index,
current_cluster)
        if not np.array_equal(clusters, previous): # als het cluster
groeide
            current_cluster += 1

    return clusters

%%timeit
clusters = DBSCAN(X)

```

6.36 ms  $\pm$  87.5  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 100 loops each)

Your code should run in less than 100ms

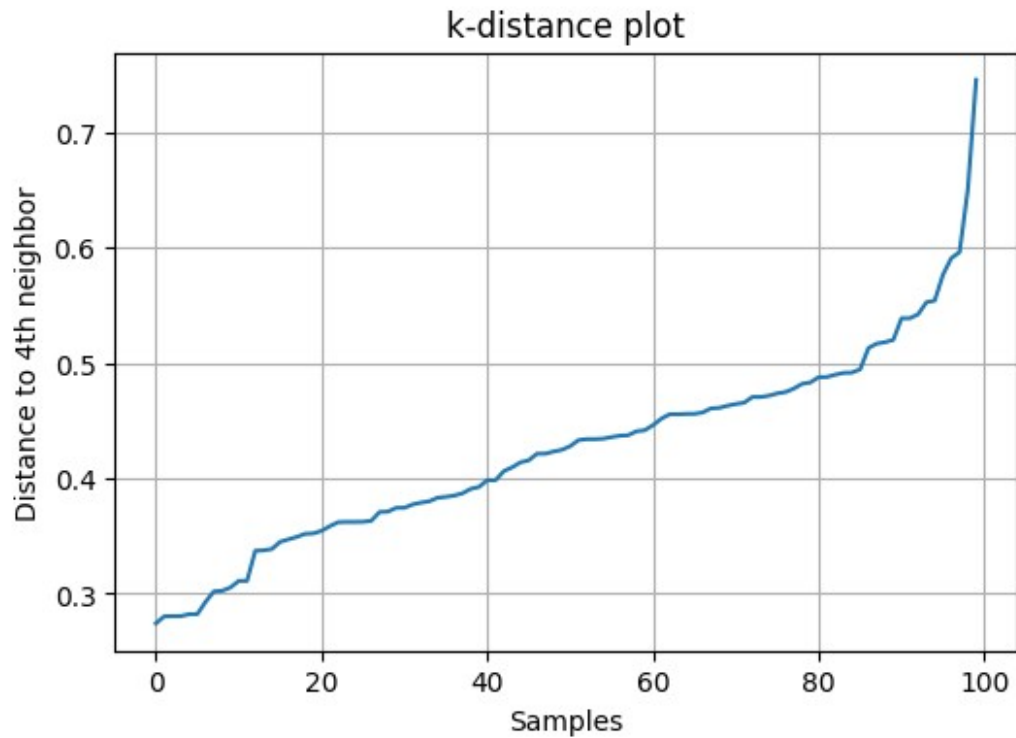
2.6 Check if your algorithm works on the following dataset. Visualise the cluster by using PCA. Make sure to first estimate the correct parameters.

```
from sklearn.datasets import make_blobs
from sklearn.decomposition import PCA
X = make_blobs(n_samples = 100, n_features = 6)[0]
X_scaled = StandardScaler().fit_transform(X)

neigh = NearestNeighbors(n_neighbors=MIN_NEIGHBOURS)
nbrs = neigh.fit(X_scaled)
distances, _ = nbrs.kneighbors(X_scaled)

# Sorteert de afstanden van de MIN_NEIGHBOURS-de buur
k_distances = np.sort(distances[:, MIN_NEIGHBOURS - 1])

plt.figure(figsize=(6, 4))
plt.plot(k_distances)
plt.title("k-distance plot")
plt.xlabel("Samples")
plt.ylabel("Distance to {}th neighbor".format(MIN_NEIGHBOURS))
plt.grid(True)
plt.show()
```



```
clusters = DBSCAN(X_scaled)
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

plt.figure(figsize=(6, 5))
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=clusters, cmap="tab10", s=30)
plt.title("Custom DBSCAN Clusters (PCA projectie)")
plt.xlabel("PCA component 1")
plt.ylabel("PCA component 2")
plt.grid(True)
plt.show()
```

